

GitHub: <https://github.com/AnishKoppula1/NeuralAssignment9.git>

## 1. Add one more hidden layer to autoencoder

```
In [6]: input_img = Input(shape=(784,))
#Adding hidden layer to encoding
hiddenLayer_en=Dense(512,activation='relu')(input_img)
# "encoded" is the encoded representation of the input
encoded = Dense(encoding_dim, activation='relu')(hiddenLayer_en) #Undercomplete Encoding
#Adding hidden layer to decoding
hiddenLayer_de=Dense(512,activation='relu')(encoded)
# "decoded" is the lossy reconstruction of the input
decoded = Dense(784, activation='sigmoid')(hiddenLayer_de)

# this model maps an input to its reconstruction
autoencoder = Model(input_img, decoded)
# this model maps an input to its encoded representation
autoencoder.compile(optimizer='adam', loss='binary_crossentropy', metrics=['accuracy'])

from keras.datasets import mnist, fashion_mnist
import numpy as np

(x_train, y_train), (x_test, y_test) = fashion_mnist.load_data()

#Converting into float & Scaling Data
x_train = x_train.astype('float32') / 255.
x_test = x_test.astype('float32') / 255.

#Setting Up data from 28*28 to 784 for the width & height of image
x_train = x_train.reshape((len(x_train), np.prod(x_train.shape[1:])))
x_test = x_test.reshape((len(x_test), np.prod(x_test.shape[1:])))

#Fitting/training the model
autoencoder.fit(x_train, x_train,
                epochs=5,
                batch_size=128,
                shuffle=True,
                validation_data=(x_test, x_test))


Epoch 1/5
469/469 ————— 8s 13ms/step - accuracy: 0.0095 - loss: 0.3772 - val_accuracy: 0.0191 - val_loss: 0.2930
Epoch 2/5
469/469 ————— 5s 11ms/step - accuracy: 0.0204 - loss: 0.2887 - val_accuracy: 0.0254 - val_loss: 0.2840
Epoch 3/5
469/469 ————— 5s 11ms/step - accuracy: 0.0243 - loss: 0.2809 - val_accuracy: 0.0268 - val_loss: 0.2801
Epoch 4/5
469/469 ————— 6s 13ms/step - accuracy: 0.0280 - loss: 0.2772 - val_accuracy: 0.0298 - val_loss: 0.2780
Epoch 5/5
469/469 ————— 6s 13ms/step - accuracy: 0.0333 - loss: 0.2753 - val_accuracy: 0.0343 - val_loss: 0.2763

Out[6]: <keras.src.callbacks.history.History at 0x2bf8350b890>
```


## 2. Do the prediction on the test data and then visualize one of the reconstructed version of that test data. Also, visualize the same test data before reconstruction using Matplotlib

```
In [7]: #predicting on the test data  
prediction = autoencoder.predict(x_test)  
313/313 ————— 1s 3ms/step
```

```
In [8]: #Input Image  
from matplotlib import pyplot as plt  
plt.imshow(x_test[50].reshape(28,28))  
plt.show()
```



```
In [9]: #reconstructed Image  
from matplotlib import pyplot as plt  
plt.imshow(prediction[50].reshape(28,28))  
plt.show()
```



### 3. Repeat the question 2 on the denoising autoencoder

```
In [10]: from keras.layers import Input, Dense
         from keras.models import Model

         # this is the size of our encoded representations
         encoding_dim = 32 # 32 floats -> compression of factor 24.5, assuming the input is 784 floats

         # this is our input placeholder
         input_img = Input(shape=(784,))
         # "encoded" is the encoded representation of the input
         encoded = Dense(encoding_dim, activation='relu')(input_img) #Undercomplete Encoding
         # "decoded" is the lossy reconstruction of the input
         decoded = Dense(784, activation='sigmoid')(encoded)
         # this model maps an input to its reconstruction
         autoencoder = Model(input_img, decoded)
         # this model maps an input to its encoded representation
         autoencoder.compile(optimizer='adam', loss='binary_crossentropy', metrics=['accuracy'])

         from keras.datasets import mnist, fashion_mnist
         import numpy as np

         (x_train, y_train), (x_test, y_test) = fashion_mnist.load_data()
         x_train = x_train[:6000]
         x_test = x_test[:1000]


         #Converting into float & Scaling Data
         x_train = x_train.astype('float32') / 255.
         x_test = x_test.astype('float32') / 255.

         #Setting Up data from 28*28 to 784 for the width & height of image
         x_train = x_train.reshape((len(x_train), np.prod(x_train.shape[1:])))
         x_test = x_test.reshape((len(x_test), np.prod(x_test.shape[1:])))


In [11]: noise_factor = 0.5
         x_train_noisy = x_train + noise_factor * np.random.normal(loc=0.0, scale=1.0, size=x_train.shape)
         x_test_noisy = x_test + noise_factor * np.random.normal(loc=0.0, scale=1.0, size=x_test.shape)

In [12]: history = autoencoder.fit(x_train_noisy, x_train,
                                   epochs=10,
                                   batch_size=256,
                                   shuffle=True,
                                   validation_data=(x_test_noisy, x_test_noisy))
```

```
In [13]: from matplotlib import pyplot as plt
plt.imshow(x_train[50].reshape(28,28))
plt.show()
```



```
In [14]: from matplotlib import pyplot as plt
plt.imshow(x_train_noisy[50].reshape(28,28))
plt.show()
```



#### 4. plot loss and accuracy using the history object

```
In [18]: M autoencoder.metrics_names
```

```
Out[18]: ['loss', 'compile_metrics']
```

```
In [19]: M import matplotlib.pyplot as plt
plt.plot(history.history['accuracy'])
plt.plot(history.history['loss'])
plt.title('model accuracy vs loss')
plt.xlabel('epoch')
plt.legend(['accuracy', 'loss'], loc='upper left')
plt.show()
```

